

```
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://apache.mirror colo-serv.net/spark/spark-2.4
!tar xf spark-2.4.7-bin-hadoop2.7.tgz
!pip install -q findspark
```

```
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64
os.environ["SPARK_HOME"] = "/content/spark-2.4.7-bin-hadoop2
```

```
import findspark
findspark.init("spark-2.4.7-bin-hadoop2.7")# SPARK_HOME
```

```
import pyspark
from pyspark.sql import *
from pyspark.sql.functions import *
from pyspark import SparkContext, SparkConf
```

```
sc = SparkContext.getOrCreate()
spark = SparkSession.builder.getOrCreate()
```

```
!wget http://files.grouplens.org/datasets/movielens/ml-lates
!unzip ml-latest-small.zip
```

```
--2021-04-09 05:53:39-- http://files.grouplens.org/datasets/movielens/ml-latest-small.zip
Resolving files.grouplens.org (files.grouplens.org)...
Connecting to files.grouplens.org (files.grouplens.org)
HTTP request sent, awaiting response... 200 OK
Length: 978202 (955K) [application/zip]
Saving to: 'ml-latest-small.zip'
```

```
ml-latest-small.zip 100%[=====>] 955.28K
```

```
2021-04-09 05:53:39 (6.24 MB/s) - 'ml-latest-small.zip'
```

```
Archive:  ml-latest-small.zip
creating: ml-latest-small/
```

```
inflating: ml-latest-small/links.csv
inflating: ml-latest-small/tags.csv
inflating: ml-latest-small/ratings.csv
inflating: ml-latest-small/README.txt
inflating: ml-latest-small/movies.csv
```

```
ratings = spark.read.csv("ml-latest-small/ratings.csv", head
movies = spark.read.csv("ml-latest-small/movies.csv", header
```

```
ratings.printSchema()
movies.printSchema()
```

```
root
|-- userId: string (nullable = true)
|-- movieId: string (nullable = true)
|-- rating: string (nullable = true)
|-- timestamp: string (nullable = true)
```

```
root
|-- movieId: string (nullable = true)
|-- title: string (nullable = true)
|-- genres: string (nullable = true)
```

```
ratings.take(3)
```

```
[Row(userId='1', movieId='1', rating='4.0', timestamp='
Row(userId='1', movieId='3', rating='4.0', timestamp='
Row(userId='1', movieId='6', rating='4.0', timestamp='
```

```
ratings.show()
```

```
+-----+-----+-----+-----+
|userId|movieId|rating|timestamp|
+-----+-----+-----+-----+
|      1|      1|   4.0|964982703|
```

1	3	4.0	964981247
1	6	4.0	964982224
1	47	5.0	964983815
1	50	5.0	964982931
1	70	3.0	964982400
1	101	5.0	964980868
1	110	4.0	964982176
1	151	5.0	964984041
1	157	5.0	964984100
1	163	5.0	964983650
1	216	5.0	964981208
1	223	3.0	964980985
1	231	5.0	964981179
1	235	4.0	964980908
1	260	5.0	964981680
1	296	3.0	964982967
1	316	3.0	964982310
1	333	5.0	964981179
1	349	4.0	964982563

+-----+-----+-----+-----+

only showing top 20 rows

`movies.show()`

☞

movieId	title	genres
1	Toy Story (1995)	Adventure Animati...
2	Jumanji (1995)	Adventure Childre...
3	Grumpier Old Men ...	Comedy Romance
4	Waiting to Exhale...	Comedy Drama Romance
5	Father of the Bri...	Comedy
6	Heat (1995)	Action Crime Thri...
7	Sabrina (1995)	Comedy Romance
8	Tom and Huck (1995)	Adventure Children
9	Sudden Death (1995)	Action
10	GoldenEye (1995)	Action Adventure ...
11	American Presiden...	Comedy Drama Romance
12	Dracula: Dead and...	Comedy Horror

13	Balto (1995)	Adventure	Animati...
14	Nixon (1995)		Drama
15	Cutthroat Island ...	Action	Adventure ...
16	Casino (1995)		Crime Drama
17	Sense and Sensibi...		Drama Romance
18	Four Rooms (1995)		Comedy
19	Ace Ventura: When...		Comedy
20	Money Train (1995)	Action	Comedy Cri...

+-----+-----+-----+
only showing top 20 rows

```
print("In total there are {0} movies".format(movies.count()))
```

In total there are 9742 movies

```
# TODO
# How many ratings are there for each movie id?
# Sort descending by the count of ratings.
# Do not use SQL for this cell.
```

```
movie_counts = ratings.groupBy("movieId")\
    .agg(count("rating")\
        .alias("cnt"))\
    .sort(desc("cnt"))
```

```
movie_counts.show()
```

```
+-----+---+
|movieId|cnt|
+-----+---+
|      356|329|
|      318|317|
|      296|307|
|      593|279|
```

	2571		278	
	260		251	
	480		238	
	110		237	
	589		224	
	527		220	
	2959		218	
	1		215	
	1196		211	
	2858		204	
	50		204	
	47		203	
	780		202	
	150		201	
	1198		200	
	4993		198	

+-----+-----+

only showing top 20 rows

```
# TODO
# Register the ratings and movies dataframes as tables
# so that we can execute sql queries on them.
ratings.registerTempTable('ratings')
movies.registerTempTable('movies')
```

```
# TODO
# Use SQL
# How many ratings are there for each movie id?
# Sort descending by the count of ratings.
query = """
select r.movieId, count(r.rating) as cnt
from ratings r
group by r.movieId
order by cnt desc
"""
result = spark.sql(query)
```

```
result = spark.sql(query)
result.show()
```

```
+-----+-----+
|movieId|cnt|
+-----+-----+
|      356|329|
|      318|317|
|      296|307|
|      593|279|
|     2571|278|
|      260|251|
|      480|238|
|      110|237|
|      589|224|
|      527|220|
|     2959|218|
|         1|215|
|     1196|211|
|     2858|204|
|        50|204|
|        47|203|
|      780|202|
|      150|201|
|     1198|200|
|     4993|198|
+-----+-----+
```

only showing top 20 rows

```
# TODO
# Use SQL
# Find the average rating for each movie that has more than
# Sort descending by average rating.
```

```
query = """
select movieId, title, avg(rating) as avgrating
from ratings natural join movies
```

```

group by movieId, title
having count(rating) >50
order by avgrating desc
"""

result = spark.sql(query)
result.show()

```

movieId	title	avgrating
318	Shawshank Redempt...	4.429022082018927
858	Godfather, The (1...	4.2890625
2959	Fight Club (1999)	4.272935779816514
1276	Cool Hand Luke (1...	4.271929824561403
750	Dr. Strangelove o...	4.268041237113402
904	Rear Window (1954)	4.261904761904762
1221	Godfather: Part I...	4.25968992248062
48516	Departed, The (2006)	4.252336448598131
1213	Goodfellas (1990)	4.25
912	Casablanca (1942)	4.24
58559	Dark Knight, The ...	4.238255033557047
50	Usual Suspects, T...	4.237745098039215
1197	Princess Bride, T...	4.232394366197183
260	Star Wars: Episod...	4.231075697211155
527	Schindler's List ...	4.225
1208	Apocalypse Now (1...	4.219626168224299
2329	American History ...	4.217054263565892
1196	Star Wars: Episod...	4.2156398104265405
1252	Chinatown (1974)	4.211864406779661
1198	Raiders of the Lo...	4.2075

only showing top 20 rows

```

movies.show()

```

movieId	title	genres
---------	-------	--------

1	Toy Story (1995)	Adventure Animati...
2	Jumanji (1995)	Adventure Childre...
3	Grumpier Old Men ...	Comedy Romance
4	Waiting to Exhale...	Comedy Drama Romance
5	Father of the Bri...	Comedy
6	Heat (1995)	Action Crime Thri...
7	Sabrina (1995)	Comedy Romance
8	Tom and Huck (1995)	Adventure Children
9	Sudden Death (1995)	Action
10	GoldenEye (1995)	Action Adventure ...
11	American Presiden...	Comedy Drama Romance
12	Dracula: Dead and...	Comedy Horror
13	Balto (1995)	Adventure Animati...
14	Nixon (1995)	Drama
15	Cutthroat Island ...	Action Adventure ...
16	Casino (1995)	Crime Drama
17	Sense and Sensibi...	Drama Romance
18	Four Rooms (1995)	Comedy
19	Ace Ventura: When...	Comedy
20	Money Train (1995)	Action Comedy Cri...

only showing top 20 rows

TODO

Here we want to extract the genres from the movies table.

This can be done easily using the underlying RDD of the mo

```
genres_rdd = movies.rdd.map(lambda row: (row.movieId, row.ge
```

```
genres_rdd.take(20)
```

```
[('1', 'Adventure|Animation|Children|Comedy|Fantasy'),
 ('2', 'Adventure|Children|Fantasy'),
 ('3', 'Comedy|Romance'),
 ('4', 'Comedy|Drama|Romance'),
 ('5', 'Comedy'),
 ('6', 'Action|Crime|Thriller'),
```



```
( '7', 'Comedy|Romance' ),
( '8', 'Adventure|Children' ),
( '9', 'Action' ),
( '10', 'Action|Adventure|Thriller' ),
( '11', 'Comedy|Drama|Romance' ),
( '12', 'Comedy|Horror' ),
( '13', 'Adventure|Animation|Children' ),
( '14', 'Drama' ),
( '15', 'Action|Adventure|Romance' ),
( '16', 'Crime|Drama' ),
( '17', 'Drama|Romance' ),
( '18', 'Comedy' ),
( '19', 'Comedy' ),
( '20', 'Action|Comedy|Crime|Drama|Thriller' )]
```

TODO

Now we want to create a flattened out RDD of (movieId, gen

Use flatMapValues on the previous RDD.

See: <https://spark.apache.org/docs/1.1.1/api/python/pyspar>

Call the result RDD: moviegenres_flat_rdd

```
def f(genres_rdd) : return genres_rdd.split('|')
```

```
moviegenres_flat_rdd=genres_rdd.flatMapValues(f).collect()
```

```
genres_rdd.flatMapValues(f).take(20)
```

```
[ ( '1', 'Adventure' ),
  ( '1', 'Animation' ),
  ( '1', 'Children' ),
  ( '1', 'Comedy' ),
  ( '1', 'Fantasy' ),
  ( '2', 'Adventure' ),
  ( '2', 'Children' ),
  ( '2', 'Fantasy' ),
  ( '3', 'Comedy' ),
  ( '3', 'Romance' ),
  ( '4', 'Comedy' ),
  ( '4', 'Drama' ),
  ( '4', 'Romance' ),
  ( '5', 'Comedy' ),
  ( '6', 'Action' ),
```

```
('6', 'Crime'),  
('6', 'Thriller'),  
('7', 'Comedy'),  
('7', 'Romance'),  
('8', 'Adventure')]
```

```
# TODO
```

```
# Now convert the last RDD into a dataframe for further anal
```

```
# Just uncomment the following lines.
```

```
moviegenres = spark.createDataFrame(moviegenres_flat_rdd).to  
moviegenres.show()
```

```
# Register dataframe as a table.
```

```
moviegenres.registerTempTable('moviegenres')
```

```
+-----+-----+  
|movieId|  genre|  
+-----+-----+  
|      1|Adventure|  
|      1|Animation|  
|      1| Children|  
|      1|  Comedy|  
|      1|  Fantasy|  
|      2|Adventure|  
|      2| Children|  
|      2|  Fantasy|  
|      3|  Comedy|  
|      3|  Romance|  
|      4|  Comedy|  
|      4|   Drama|  
|      4|  Romance|  
|      5|  Comedy|  
|      6|  Action|  
|      6|   Crime|  
|      6|Thriller|  
|      7|  Comedy|  
|      7|  Romance|  
|      8|Adventure|
```

```
+-----+-----+
only showing top 20 rows
```

```
# TODO
# Use SQL
# Find the average rating for each genre that has at least 5
# Order descending by average rating.
query = """
select genre, avg(rating) as avgrating
from ratings natural join moviegenres
group by genre
having count(rating) >50
order by avgrating desc
"""

result = spark.sql(query)
result.show()
```

```
+-----+-----+
|      genre|      avgrating|
+-----+-----+
| Film-Noir| 3.920114942528736|
|      War| 3.8082938876312|
| Documentary| 3.797785069729286|
|      Crime| 3.658293867274144|
|      Drama| 3.6561844113718758|
|      Mystery| 3.632460255407871|
| Animation| 3.6299370349170004|
|      IMAX| 3.618335343787696|
| Western| 3.583937823834197|
| Musical| 3.5636781053649105|
| Adventure| 3.5086089151939075|
| Romance| 3.5065107040388437|
| Thriller| 3.4937055799183425|
| Fantasy| 3.4910005070136894|
| Sci-Fi| 3.455721162210752|
| Action| 3.447984331646809|
```

	Children		3.412956125108601	
	Comedy		3.3847207640898267	
	Horror		3.258195034974626	
+	-----	+	-----	+

```
# Uncomment following lines.
result_pd = result.toPandas()
result_pd.head()
```

	genre	avgrating
0	Film-Noir	3.920115
1	War	3.808294
2	Documentary	3.797785
3	Crime	3.658294
4	Drama	3.656184

```
# Uncomment following line.
result_pd.plot(kind="bar", x="genre", y="avgrating")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f15d162eb9

